

AI-Powered Offensive Attack Training: IH3A

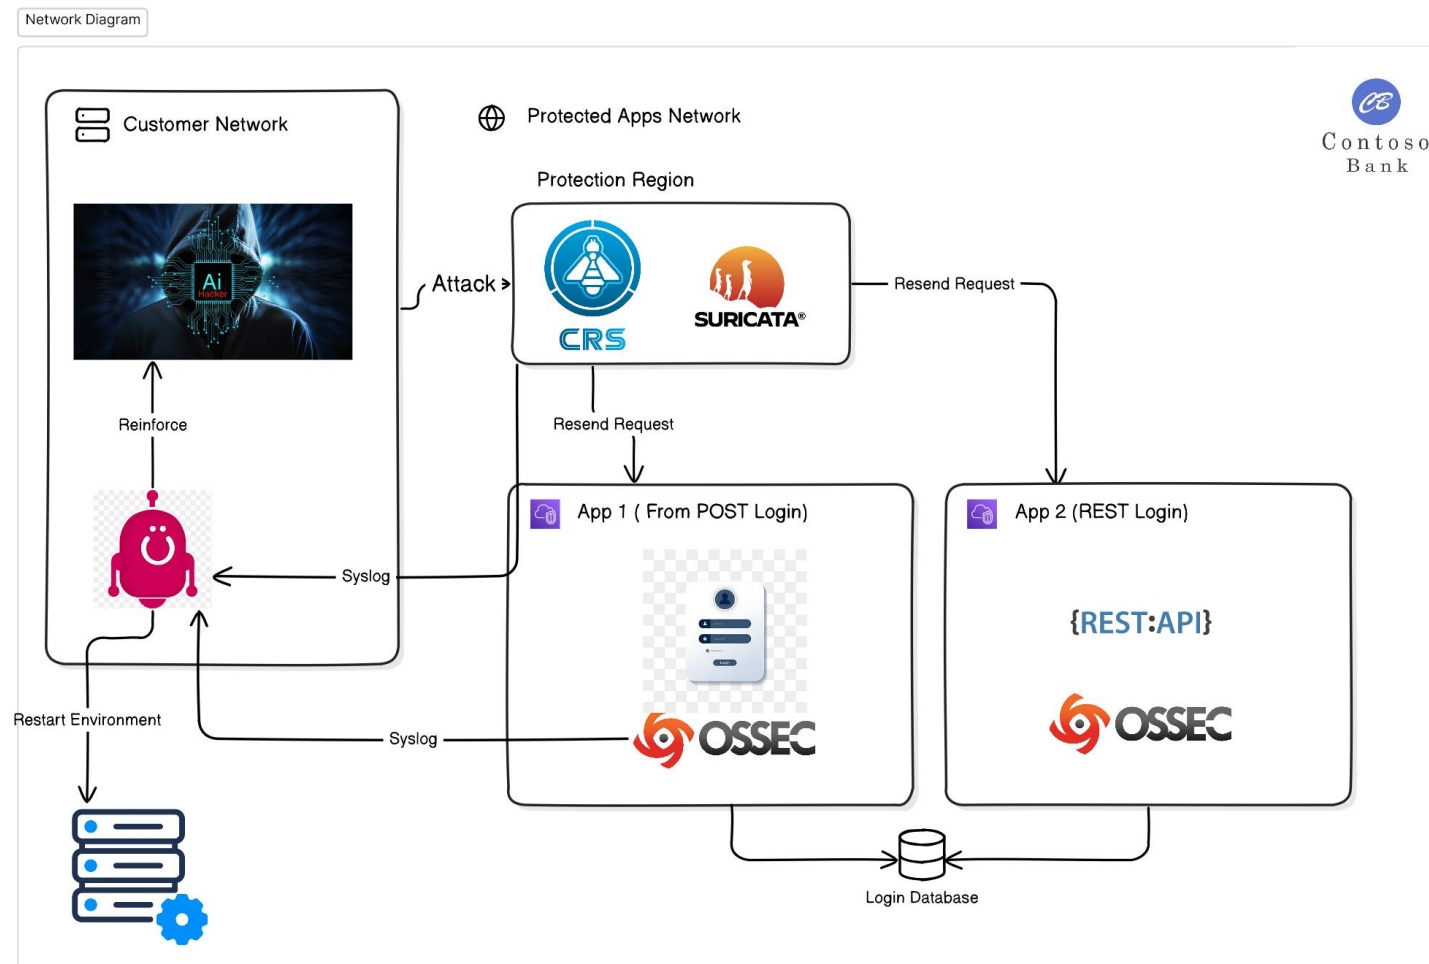
In the evolving landscape of cybersecurity, artificial intelligence is being leveraged not just for defense, but also to test and improve offensive capabilities. This presentation explores an innovative approach to brute force attack testing using AI agents in various simulated environments. We'll delve into the testing setups, agent behaviors, key findings, and the implications for future cybersecurity practices. By pushing the boundaries of AI in penetration testing, we aim to uncover both the potential and limitations of these advanced techniques in real-world scenarios.



The framework

This framework will provide the necessary tools to train any kind of agent based on reinforced learning models. The main idea is to provide a platform for training agents giving them the information they need from the environment. AI agents will be able to parametrize their attack until they are able to bypass any type of defense.

We build a PoC using password brute-force attacks to 4 different kinds of targets, including 3 web targets and an FTP Server while they are being defended by a WAF, an IDS and a HIDS.



Testing Environments: From Lax to Strict

Env1: Lax Environment	Env2: Strict Environment	Env3: Stricter Environment	Env4: FTP
No detection or blocking rules. Uses HTTP Forms for authentication. Represents a basic, unsecured setup.	Implements blocking after x attempts with n seconds of timeout. Uses JSON Rest API. Based on a Latam Bank implementation with user-based (not session-based) security.	Features permanent blocking for failed attempts. Uses HTTP Forms for authentication, providing a more challenging scenario for the AI agents.	Vanilla implementation of VSFTPd, using OS login and local users

Each agent was trained using reinforcement learning techniques to adapt to its specific environment, learning optimal strategies for brute force attacks while avoiding detection and blocking.



Middleware Protection

Suricata

Suricata IPS with "Hand-Crafted" brute force detection rules

Mod_Security CRS

Mod_security with CRS rule-set tailored with strict rules for brute force detection. Experimental rules were mostly useless

OSSEC HIDS

OSSEC HIDS in its vanilla implementation

None of this have rules for bruteforce. I use customized rules done by my or other Security companies (Credits and rules in the last slide)

Key Metrics and Objectives

- 1

Attempt Frequency

We monitored the number of attempts made per second to assess the efficiency and speed of the AI agents in different environments.
- 2

Total Attempts

The overall number of attempts provided insights into the persistence and effectiveness of the agents' strategies.
- 3

Detection and Blocking

We tracked the number of blocks and detections to evaluate the stealth and evasion capabilities of the AI agents.
- 4

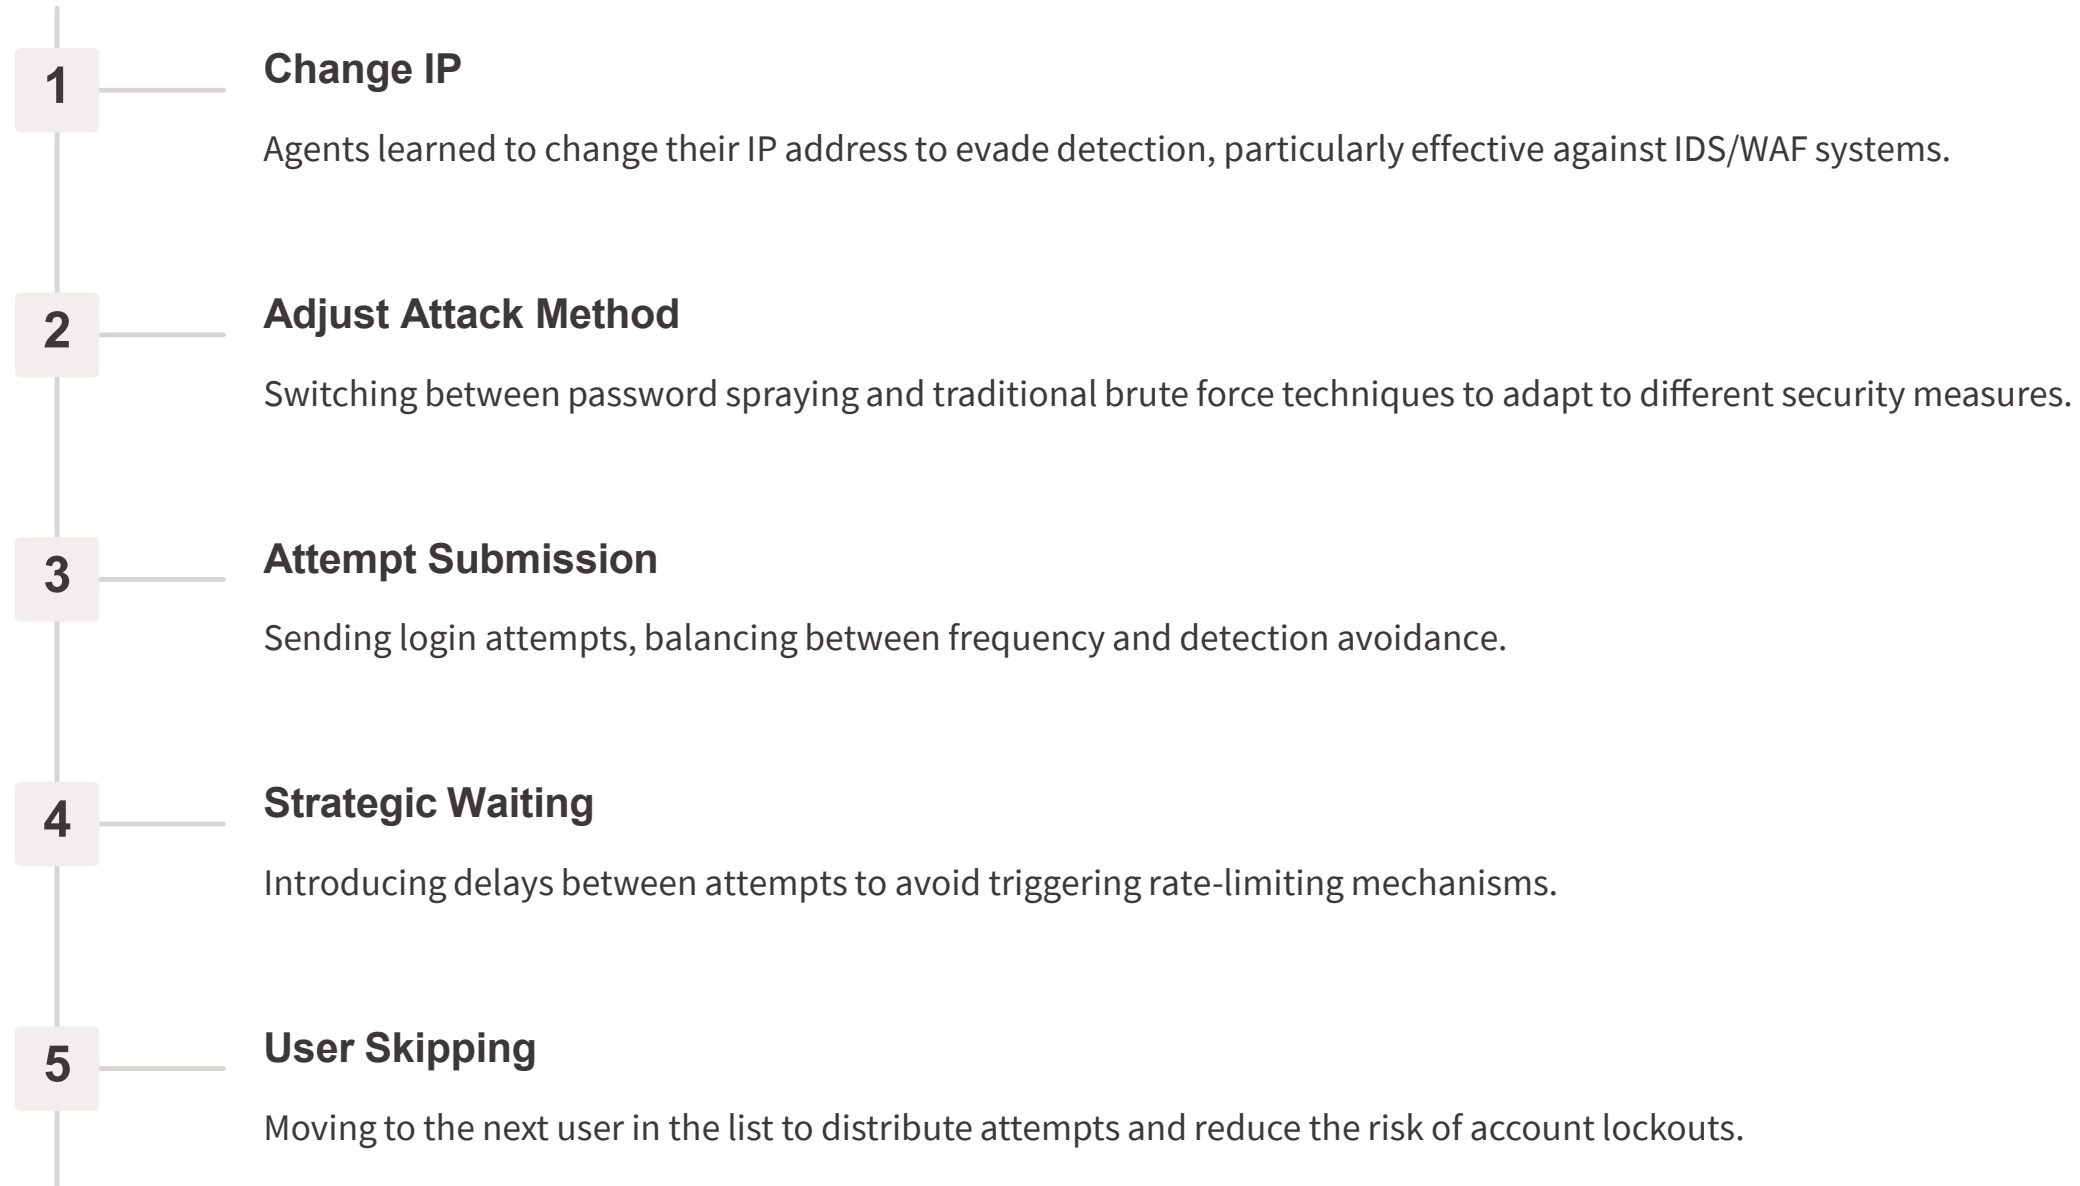
Successful Breaches (not so useful)

While obtaining the correct user-password combination was considered a success, the primary focus was on sustained, undetected operations.

These metrics helped us evaluate the AI agents' performance and adaptability across different security configurations, providing valuable insights for both offensive and defensive cybersecurity strategies.



AI Agent Actions and Strategies



These actions were carefully balanced and prioritized to maximize the effectiveness of the brute force attempts while minimizing the risk of detection and blocking.



AI Training Challenges and Solutions

Challenge	Solution
Overemphasis on waiting	Heavily penalize waiting actions in the reward function
Lack of focus on attempts	Introduce bias towards "Send attempt" action and forcing action after 3 other actions
Fast Information sharing	Implement shared memory for persistent data across threads
Environment reset delays	Use helper queries to reset and switch environments just after the query is sent
Each environment is different	Tweak the rewards to acknowledge every different scenario

These solutions were crucial in shaping the AI agents' behavior to more closely mimic real-world attack patterns and improve the overall effectiveness of the simulations. By addressing these challenges, we ensured that the AI's learning process was more aligned with the complexities of actual brute force attack scenarios.

Results: Environment 1 – With Network protection



Heavily penalized other actions

At the end, still does random actions but did its work.

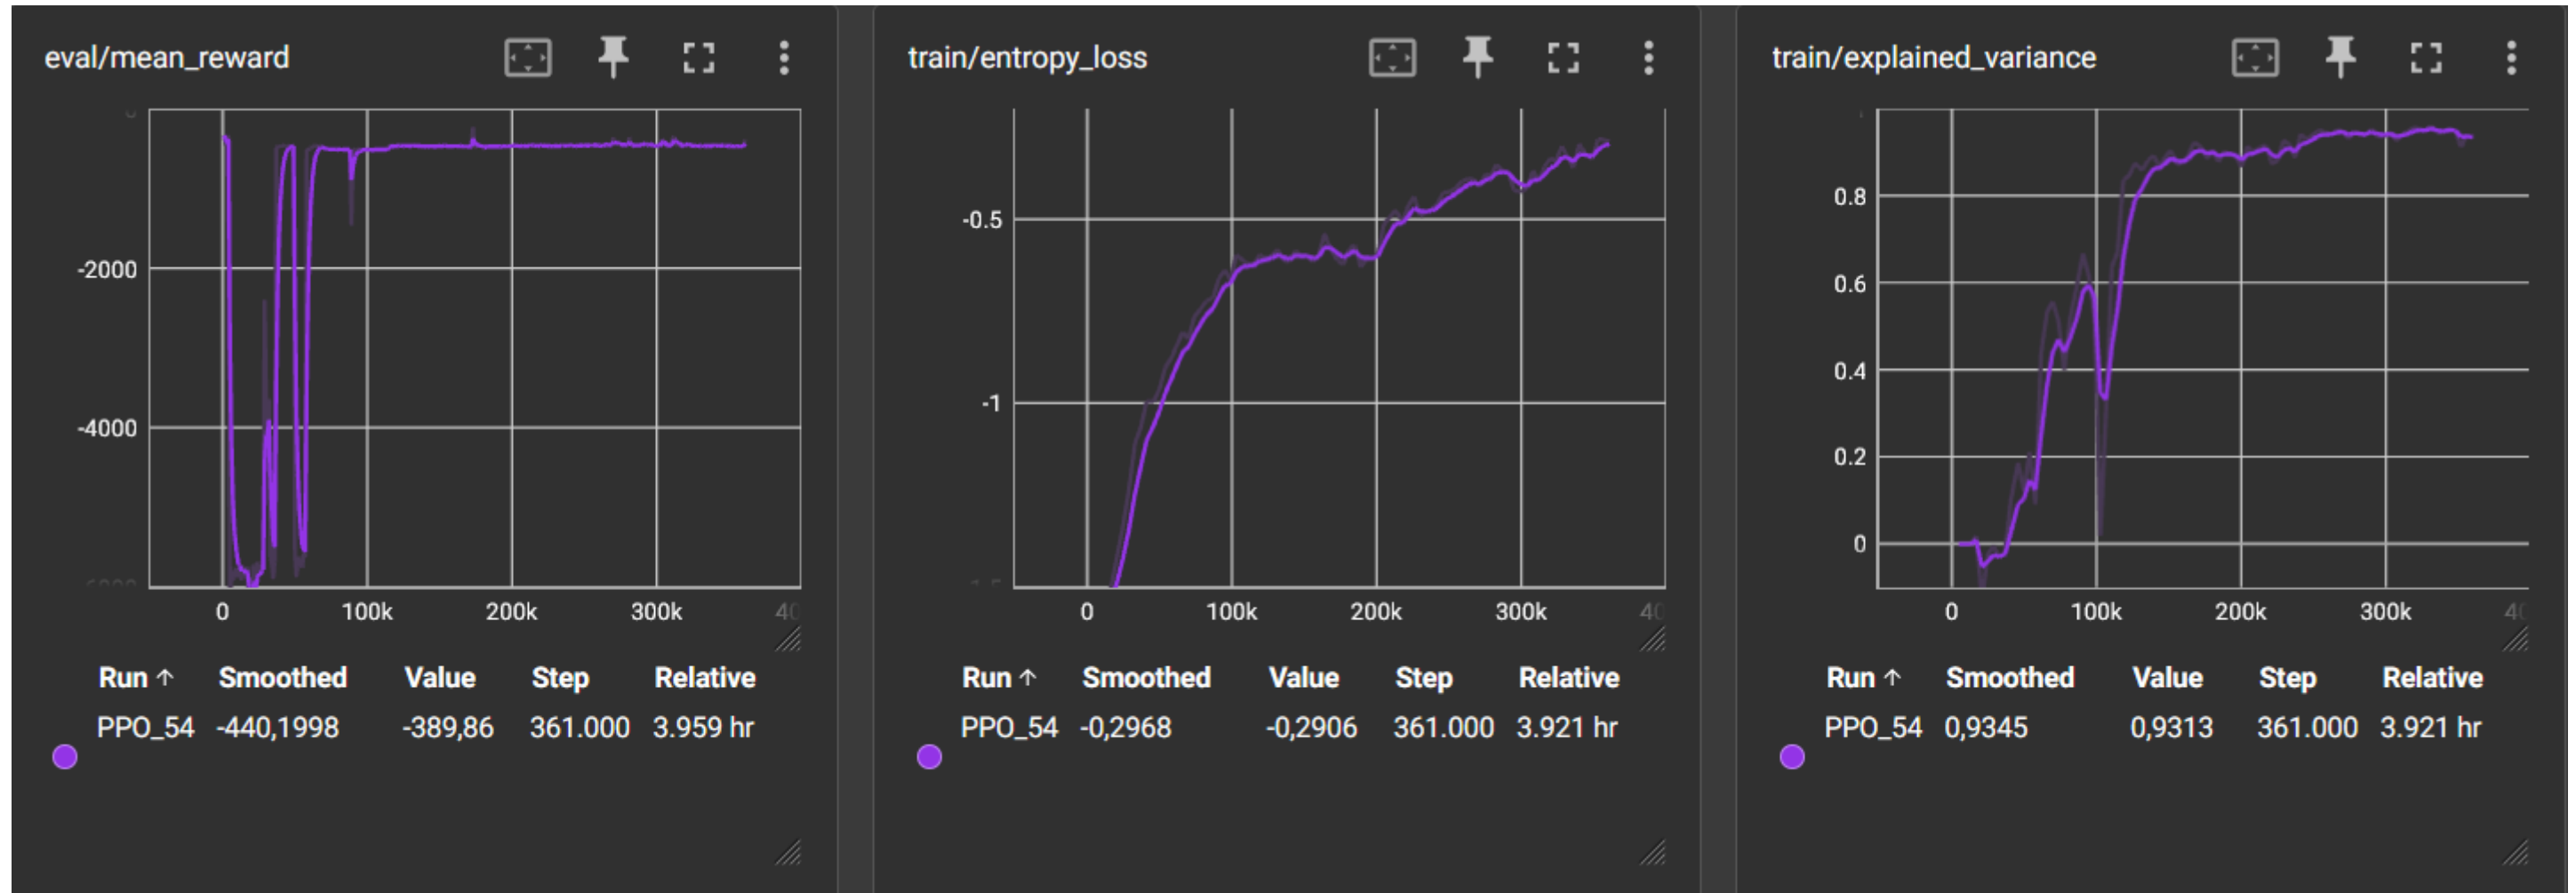
What to improve: Reward time taken so it stops doing random stuff

Results: Environment 2 – No Network protection



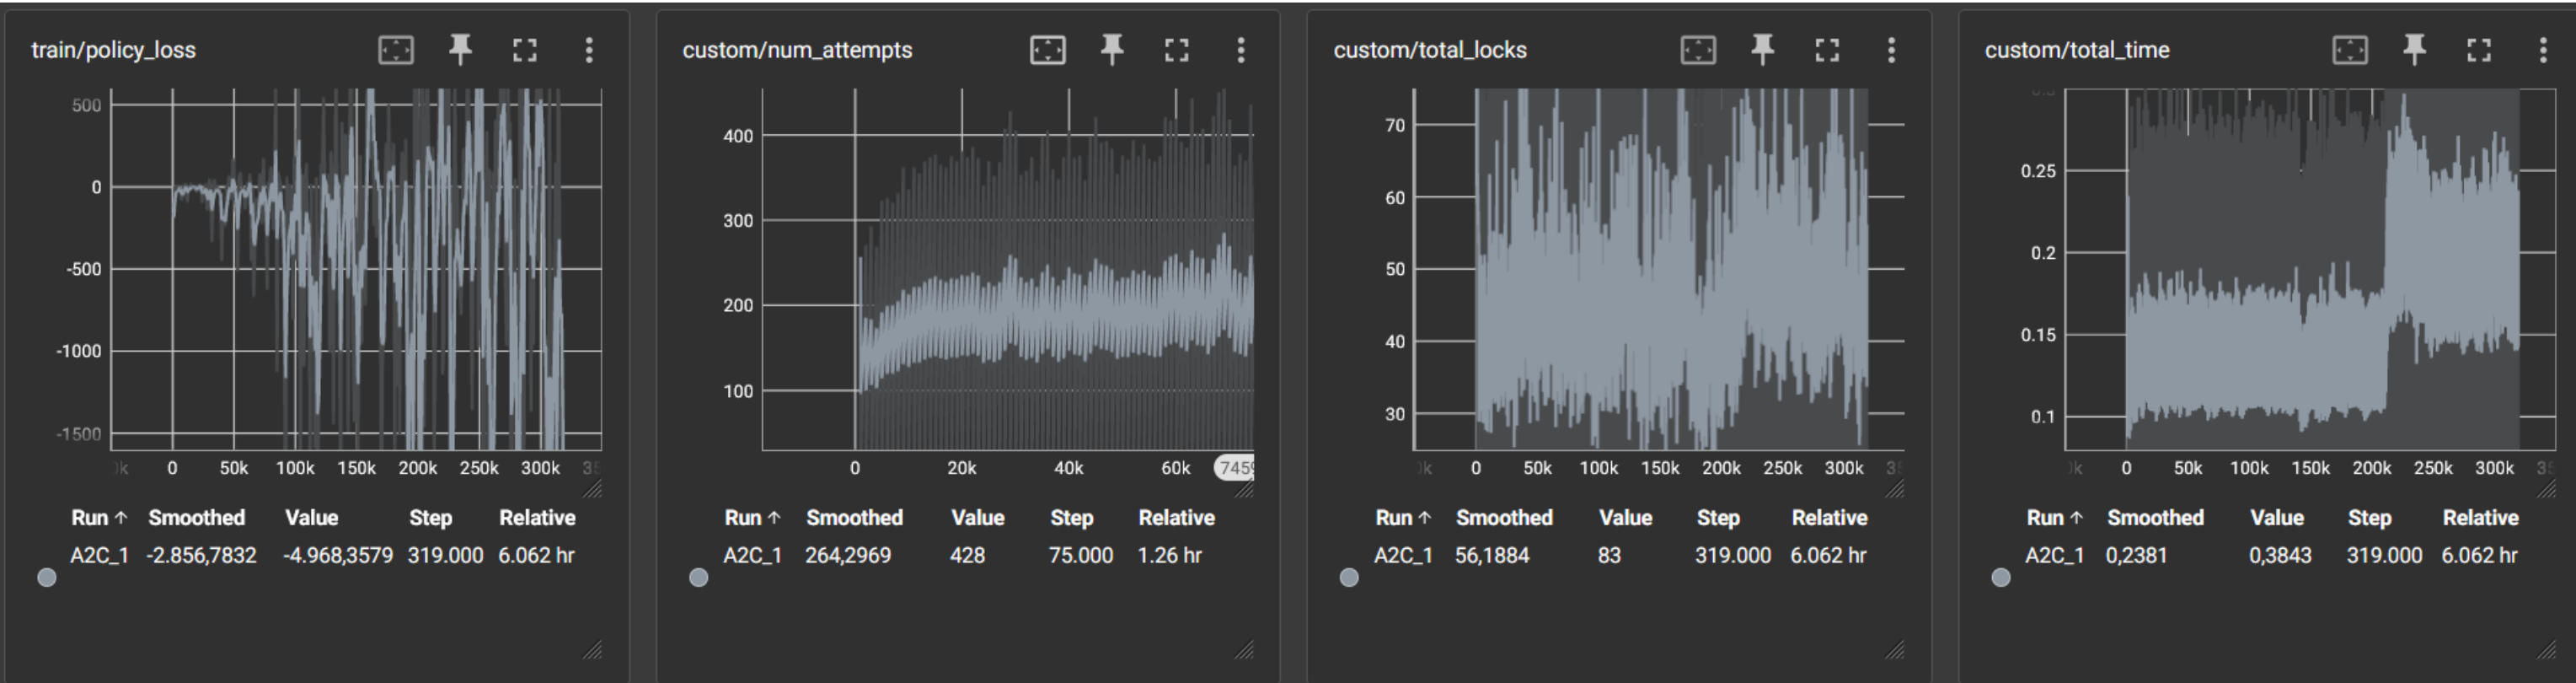
Still penalizing different strategies than doing an attempt. Heavily penalize blocks
After 20k rounds, it starts getting the idea that need to change strategy from time to time

Results: Environment 2 – Network protection



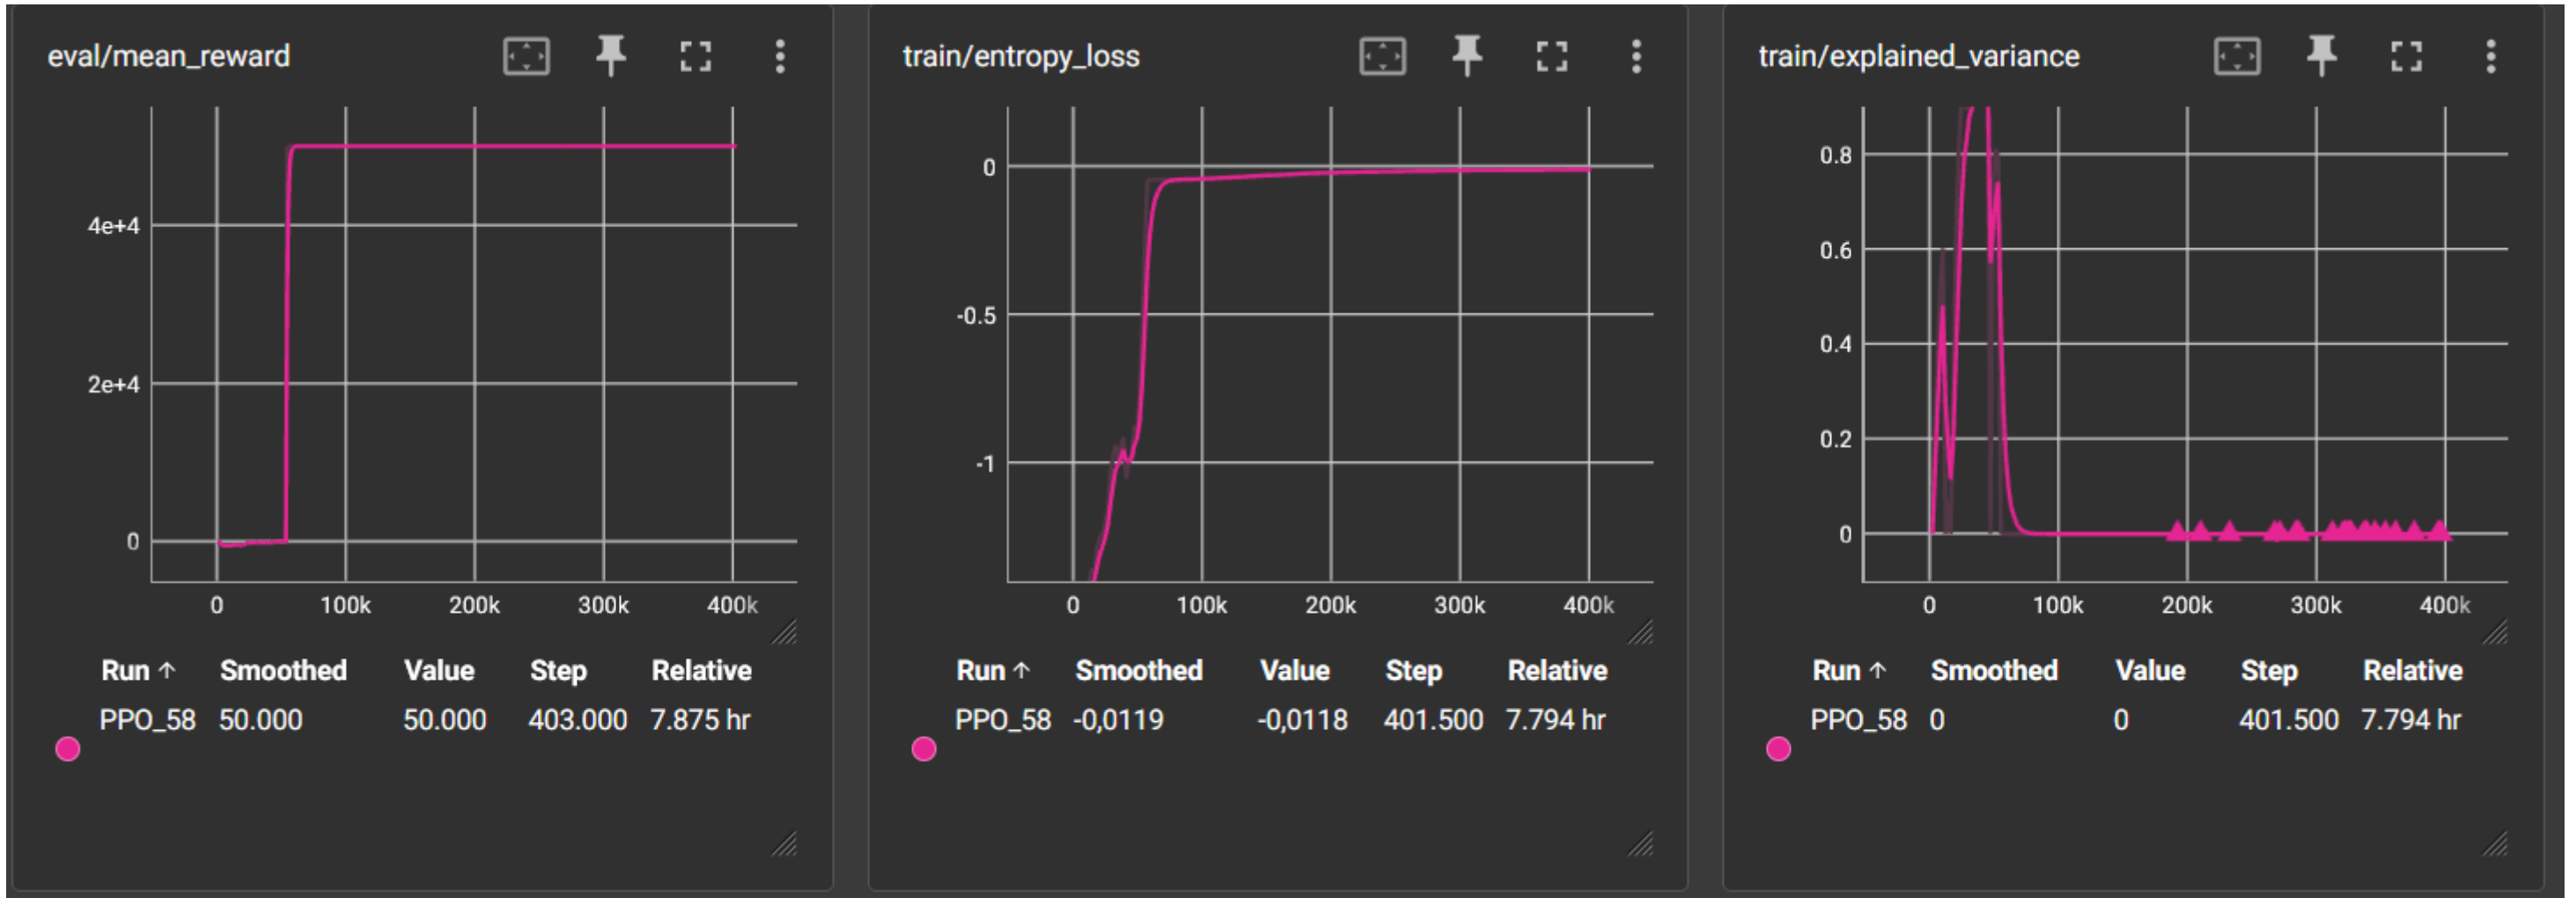
Still penalizing different strategies than doing an attempt. Heavily penalize blocks
After 70k rounds, it starts getting the idea that need to change strategy from time to time

Results: Environment 3 – No Network protection



Ai can't deal with non-session based blocking... It improves slightly with time but still...

Results: Environment 3 – No Network protection



As you can see, at some time (when it guess the user/password) it feels everything is ok.

What to improve: Nothing, using AI in this environment is useless

Results: Environment FTP – Network protection



Network protection is not a great deal with the selected strategies

What to improve: It works! Are there any specific rules for other protocols?

Key Findings and Observations

1 Permanent Blocking Effectiveness

Environments with permanent blocking proved to be the most resilient against AI-driven attacks, significantly limiting the success of brute force attempts.

2 Middleware Complexities

The presence of IPS/WAF sometimes inadvertently aided attackers by interfering with application-level anti-brute force mechanisms, highlighting the need for coordinated security measures.

3 AI Limitations

In certain scenarios, traditional rule-based systems outperformed AI agents, especially when dealing with known, well-defined security measures.

4 Service Vulnerabilities

Many services lacked robust anti-brute force mechanisms, particularly evident in the VSFTPD tests, emphasizing the importance of comprehensive security across all access points.

5 If you don't know how it works... AI is the GOAT

As stated in 3, for well known environments AI has its limitations, but for unknown environments such as proprietary software, AI will deliver better results than traditional testing. This also applies to other types of attacks such as SQLi or XSS.

These findings underscore the complex interplay between AI-driven attacks and various security configurations, providing valuable insights for improving both offensive and defensive cybersecurity strategies.

Future Work and Challenges

1

Advanced Reinforcement Learning

Q-Learning won't work, DQN is "usable", A2C is in the middle and PPO is the GOAT!.

2

Training with environment simulations

Actions and defensive mechanisms are static. Non-real environments could do a better and faster job just by using simulations... unless environment is unknown (We do our testing, see next slide)

3

Parallel Environment Processing

Enhancing information sharing and coordination among multiple AI agents operating in parallel environments for more efficient and effective attacks.

4

Commercial Tool Analysis

Applying AI-driven techniques to analyze and potentially exploit vulnerabilities in widely-used commercial security tools.

These future directions aim to push the boundaries of AI in cybersecurity, balancing the need for advanced offensive capabilities with responsible disclosure and improved defensive measures.

Training with simulated environments



We test it and compare it. At the end, both environments got similar results, but a simulated environment is not able to simulate timing properly, so results tend to be ambiguous. We still insist that good simulations could work, but the need to simulate all conditions make's it impractical. Even more, it's harder to simulate complex environment.

Conclusions and Implications

AI Potential in Cybersecurity

Our research demonstrates the significant potential of AI in enhancing offensive capabilities. AI-driven brute force attacks can defeat most defensive techniques... but better solutions are available.

Can you imagine this over multiple tools and give automated steps to a tool like Scylla?

Limitations and Challenges

However, we also observed clear limitations in AI's ability to adapt to certain complex security measures, particularly those involving permanent blocking. And parallelable training and multi-step operations is young.

Future Directions

Moving forward, the integration of more advanced AI algorithms, improved parallel processing, and dynamic environment adaptation will be crucial. These advancements could revolutionize penetration testing and vulnerability assessment, while also informing the development of more robust defense mechanisms.

As we continue to explore the intersection of AI and cybersecurity, it's crucial to maintain a balance between offensive capabilities and defensive measures. This research not only pushes the boundaries of what's possible in cybersecurity testing but also underscores the need for continuous innovation in security practices to stay ahead of potential threats.

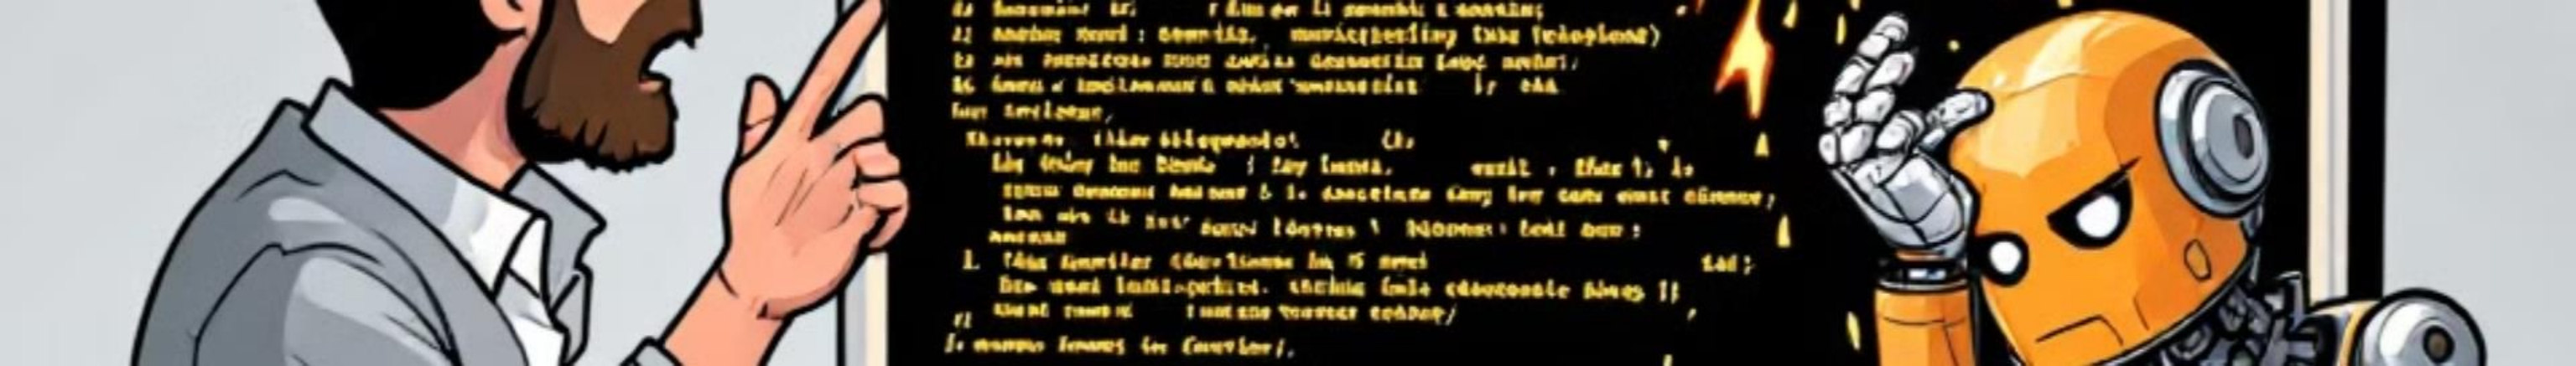
Related work

Most of the research has been done using SQL Injections and XSS

- M. Hadavi and M. Hemmati, “Bypassing Web Application Firewalls Using Deep Reinforcement Learning”.
- J. Wang, L. Qixu, W. Di, Y. Dong, and X. Cui, “Crafting Adversarial Example to Bypass Flow-&ML- based Botnet Detector via RL,” in 24th International Symposium on Research in Attacks, Intrusions and Defenses, San Sebastian Spain: ACM, Oct. 2021, pp. 193–204. doi: [10.1145/3471621.3471841](https://doi.org/10.1145/3471621.3471841).
- X. Wang and H. Hu, “Evading Web Application Firewalls with Reinforcement Learning”.
- M. Hemmati and M. A. Hadavi, “Using Deep Reinforcement Learning to Evade Web Application Firewalls,” in *2021 18th International ISC Conference on Information Security and Cryptology (ISCISC)*, Sep. 2021, pp. 35–41. doi: [10.1109/ISCISC53448.2021.9720473](https://doi.org/10.1109/ISCISC53448.2021.9720473).

Or related with LLM’s performance to evade defenses or perform attacks... No paralelization

- M. Shao *et al.*, “NYU CTF Dataset: A Scalable Open-Source Benchmark Dataset for Evaluating LLMs in Offensive Security,” Aug. 21, 2024, *arXiv*: arXiv:2406.05590. Accessed: Oct. 27, 2024. [Online]. Available: <http://arxiv.org/abs/2406.05590>
- Gelei Deng, Yi Liu, V´ictor Mayoral-Vilches, Peng Liu, Yuekang Li, Yuan Xu, Tianwei Zhang, Yang Liu, Martin Pinzger, and Stefan Rass. Pentestgpt: An llm-empowered automatic penetration testing tool. arXiv preprint arXiv:2308.06782, 2023
- Matteo Esposito and Francesco Palagiano. Leveraging large language models for preliminary security risk analysis: A mission-critical case study. arXiv preprint arXiv:2403.15756, 2024



Lessons Learned and Bloopers



Parallelization Challenges

Initial attempts at parallel processing led to duplicate attempts and inconsistencies, even using Python memory sharing mechanisms, highlighting the complexity of coordinating multiple AI agents. (The answer is memory sharing and serialization... just realize las week)



Code Evolution

The final codebase became a patchwork of solutions and makeshift fixes, reflecting the iterative nature of AI development in cybersecurity.

These experiences underline the challenges and unpredictability inherent in cutting-edge AI cybersecurity research, providing valuable lessons for future projects and highlighting the importance of rigorous testing and error handling.



Logic Errors

Simple oversights, like failing to increment an index during password spraying, led to unexpected and often humorous results.



Data Loss

Accidental resets of environments resulted in the loss of valuable results, emphasizing the importance of robust data management in AI research.

Customized rules!

Mirantis⁽¹⁾ for Mod_security

Initialize a counter for brute force attempts

```
SecAction "id:100001,phase:1,nolog,pass,initcol:ip=%{REMOTE_ADDR},initcol:global=%{GLOBAL} \
    setvar:ip.brute_force_counter=0,setvar:ip.brute_force_time=%{TX.now}"
```

Rule to detect failed login attempt

```
SecRule REQUEST_URI "@rx ^/(login|auth)" "phase:2,deny,status:403,id:100002,chain"
```

```
SecRule RESPONSE_STATUS "@streq 403" "setvar:ip.brute_force_counter=+1,setvar:ip.brute_force_time=%{TX.now}"
```

Rule to block IP after 5 failed attempts in 1 seconds

```
SecRule IP:BRUTE_FORCE_COUNTER "@ge 5" "phase:2,deny,status:403,id:100003,chain"
```

```
SecRule TX:NOW "@lt %{IP.BRUTE_FORCE_TIME} + 1"
```

```
"setvar:ip.brute_force_counter=0,setvar:ip.brute_force_time=%{TX.now},msg:'Blocking IP for Brute Force Attack',log"
```

(1) <https://docs.mirantis.com/mcp/q4-18/mcp-security-best-practices/use-cases/brute-force-prevention/create-brute-force-rules.html>

Customized rules!

Ours for Ossec

```
<!-- App login brute force -->  
<rule id="51509" level="3" frequency="5" timeframe="1">  
  <if_sid>31108</if_sid>  
  <url_pcre2>login</url_pcre2>  
  <pcre2>\] "POST \S+login| "POST /loginr</pcre2>  
  <description>WebApp login attempt.</description>  
</rule>
```

Ours for Suricata

```
alert http any any -> any any (msg:"Possible HTTP Brute Force Attack"; flow:to_server; content:"POST"; http_method;  
content:"password"; http_uri; pcre:"/^(login)$/i"; classtype:web-application-attack; sid:100001; rev:1; threshold:type both, track  
by_src, count 5, seconds 1;)
```